

IT628: Systems Programming

Process termination, wait/waitpid, exec

When does a process finish?

- A process terminates for one of 3 reasons:
 - It calls exit();
 - It returns (an int) from main
 - It receives a signal (from the OS or another process) whose default action is to terminate

- Key observation: the dying process *produces status information.*
 - Who looks at this? The parent process!

■ **void exit(int status);**

- Terminates a process with a specified status
- By convention, status of 0 is normal exit, non-zero indicates an error of some kind

```
void foo() {  
    exit(1); /* no return */  
}  
  
int main() {  
    foo(); /* no return */  
    return 0;  
}
```

Reaping Children

- **wait(): parents reap their dead children**
 - Given info about why child died, exit status, etc.

- **Two variants**
 - wait(): wait for and reap next child to exit
 - waitpid(): wait for and reap specific child

```
pid_t wait(int *stat_loc);
```

when called by a process with ≥ 1 children:

- *waits* (if needed) for a child to terminate
- *reaps* a terminated child (if ≥ 1 terminated children, arbitrarily pick one)
- *returns* reaped child's pid and exit status info via pointer (if non-NULL)

when called by a process with no children:

- return -1 immediately

```
int main() {
    pid_t cpid;
    if (fork()== 0)
        exit(0); /* terminate child */
    else
        cpid = wait(NULL); /* reaping parent */

    printf("Parent pid = %d\n", getpid());
    printf("Child pid = %d\n", cpid);

    while (1); /* Infinite loop */
}
```

```
int main() {
    if (fork()== 0){
        printf("HC: hello from child\n");
    } else {
        printf("HP: hello from parent\n");
        wait(NULL);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
}
```



```

int main() {
    if (fork()== 0){
        printf("HC: hello from child\n");
    } else {
        printf("HP: hello from parent\n");
        wait(NULL);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
}

```

A.

HP
CT
HC
Bye
Bye

B.

HP
HC
CT
Bye
Bye

C.

HP
HC
Bye
CT
Bye

D.

HC
Bye
HP
CT
Bye

E.

HC
HP
Bye
CT
Bye

```
void wait4() {  
    int stat;  
    if (fork() == 0)  
        exit(1);  
    else  
        wait(&stat);  
    printf("%d\n", stat);  
}
```

```
linux> ./wait4  
256
```

Child status information

- **status information about the child reported by wait is more than just the exit status of the child**
 - normal/abnormal termination
 - termination cause
 - exit status

WIF... macros

- **WIFEXITED (status) : child exited normally**
 - **WEXITSTATUS (status) : return code when child exits**
- **WIFSIGNALED (status) : child exited because a signal was not caught**
 - **WTERMSIG (status) : gives the number of the terminating signal**
- **WIFSTOPPED (status) : child is stopped**
 - **WSTOPSIG (status) : gives the number of the stop signal**

```
/* prints information about a signal */
```

- **void psignal(unsigned sig, const char *s);**

```
void wait5() {  
    int stat;  
    if (fork() == 0)  
        exit(1);  
    else  
        wait(&stat);  
    if (WIFEXITED(stat))  
        printf("Exit status: %d\n", WEXITSTATUS(stat));  
    else if (WIFSIGNALED(stat))  
        psignal(WTERMSIG(stat), "Exit signal");  
}
```

```
linux> ./wait5  
Exit status: 1
```

```
void wait6() {  
    int stat;  
    if (fork() == 0)  
        *(int *)NULL = 0;  
    else  
        wait(&stat);  
    if (WIFEXITED(stat))  
        printf("Exit status: %d\n", WEXITSTATUS(stat));  
    else if (WIFSIGNALED(stat))  
        psignal(WTERMSIG(stat), "Exit signal");  
    return 0;  
}
```

```
linux> ./wait6  
Exit signal: Segmentation fault
```

■ If multiple children completed, will reap in arbitrary order

```
void wait7() {  
    int i, stat;  
    pid_t pid[5];  
    for (i=0; i<5; i++)  
        if ((pid[i] = fork()) == 0) {  
            sleep(1);  
            exit(100+i);  
        }  
    for (i=0; i<5; i++) {  
        pid_t cpid = wait(&stat);  
        if (WIFEXITED(stat))  
            printf("Child %d terminated with status: %d\n",  
                   cpid, WEXITSTATUS(stat));  
    }  
}
```

waitpid(): waiting for a specific process

- Useful when parent has more than one child, or you want to check for exited child but not block

```
pid_t result =
```



The child to wait for/check on
-1 means any child

```
    waitpid(child_pid,  
            &status,  
            options);
```



0 = no options, wait until child exits
WNOHANG = don't wait, just check

- **Return value**

- pid of child, if child has exited
- 0, if using WNOHANG and child hasn't exited

■ Can use waitpid() to reap in order

```
void wait8() {  
    int i, stat;  
    pid_t pid[5];  
    for (i=0; i<5; i++)  
        if ((pid[i] = fork()) == 0) {  
            sleep(1);  
            exit(100+i);  
        }  
    for (i=0; i<5; i++) {  
        pid_t cpid = waitpid(pid[i], &stat, 0);  
        if (WIFEXITED(stat))  
            printf("Child %d terminated with status: %d\n",  
                   cpid, WEXITSTATUS(stat));  
    }  
}
```

■ Can use WNOHANG to avoid busy waiting

```
void wait9() {  
    int i, stat;  
    pid_t cpid;  
    if (fork() == 0) {  
        printf("Child pid = %d\n", getpid());  
        sleep(3);  
        exit(1);  
    } else {  
        /* use with -1 to wait on any child (with options) */  
        while ((cpid = waitpid(-1, &stat, WNOHANG)) == 0) {  
            sleep(1);  
            printf("No terminated children\n");  
        }  
        printf("Reaped %d with exit status: %d\n",  
               cpid, WEXITSTATUS(stat));  
    }  
}
```

exec(): Loading and Running Programs

- After fork, the child process is an identical duplicate of the parent process
- How do we start a new program, instead of copying the parent?
 - Use the exec() system call

```
int execve( char *filename, char *argv[ ],  
char *envp );
```

- **filename: name of the executable file to run**
- **argv: command line arguments**
- **envp: environment variable settings (e.g. \$PATH, \$HOME, etc.)**
- **returns -1 if error, otherwise doesn't return**

```
int main() { /* exec1.c */  
    char *args[2];  
    args[0] = "/bin/echo";  
    args[1] = NULL;  
    execv("/bin/echo", args);  
    return 0;  
}
```

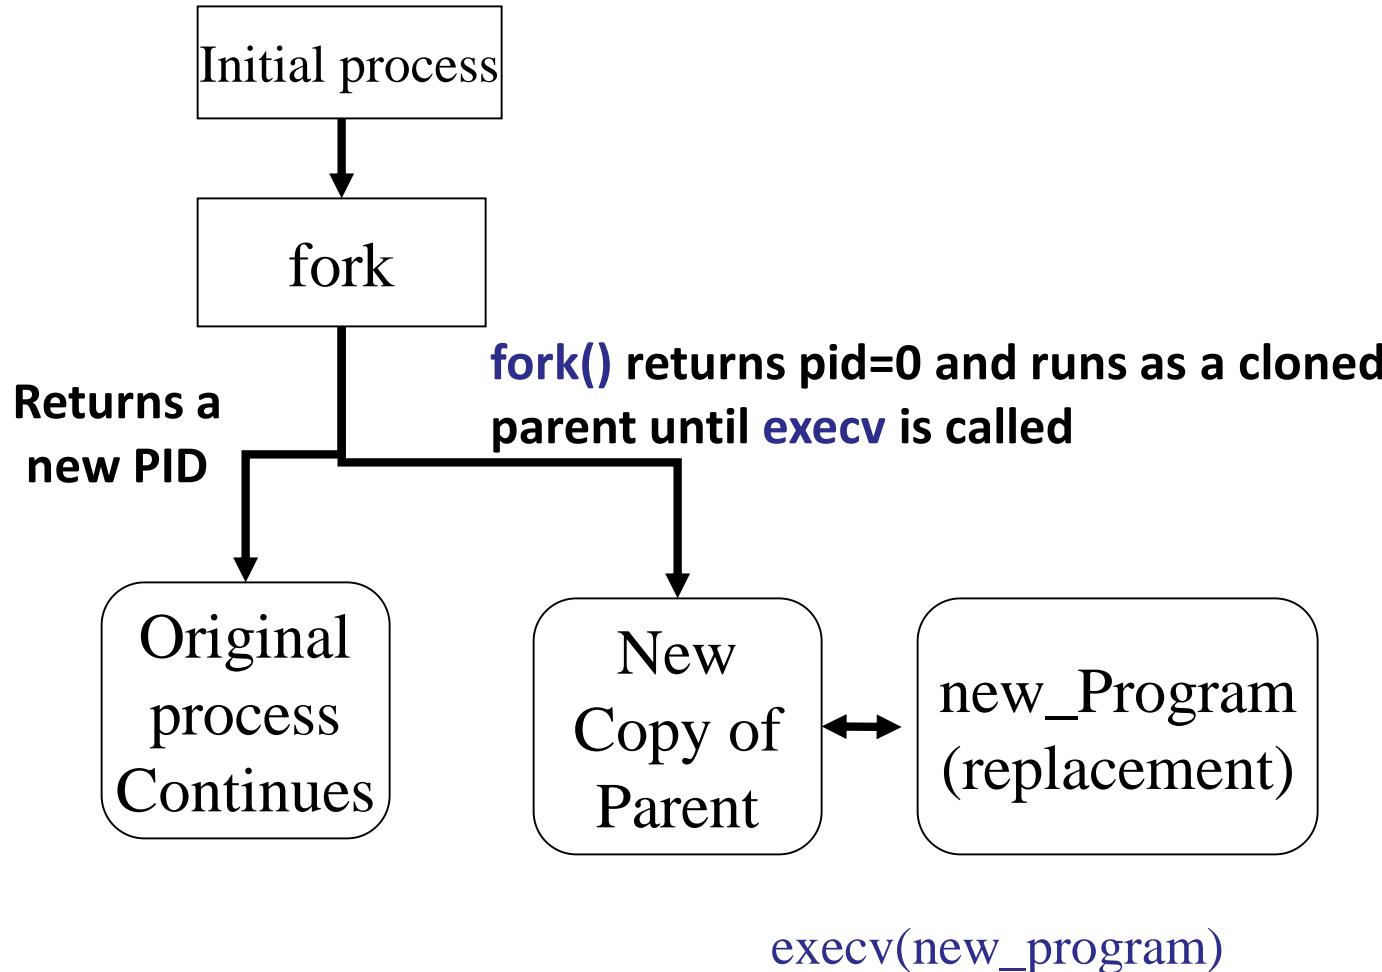
```
int main() {          /* exec2.c */
    char *args[2];
    args[0] = "/bin/echo";
    args[1] = NULL;
    printf("About to exec from process %d\n", getpid());
    sleep(1);
    execv("./exec2", args);
    printf("Done exec-ing ...\\n");
    return 0;
}
```

■ **exec() does not create a new process!**

- Replaces the address space and CPU state of the current process
 - Loads the address space from the executable file and starts it from main()
- On success, exec does not return!

■ **UNIX shells use fork-then-exec to run programs**

`execv(new_program, argv[])`



Exercise

Write a program that creates a child process, the child executes /bin/ls, and then the parent prints “done”. Make sure the word “done” is printed after the output of ls.

```
int main() {          /* exec3.c */
    if (fork() == 0) { /* Child process */
        char *args[2];
        args[0] = "/bin/ls"; /* Not required!! */
        args[1] = NULL; /* Indicate end of args array */
        execv("/bin/ls", args);
        exit(0);          /* in case exec fails! */
    }
    wait(NULL);
    printf("Done\n");
    return 0;
}
```

Waiting for a Child Process

- If a process (the parent) calls `fork()` to create a process (the child), the parent doesn't automatically wait for the child to finish. The parent must call `wait`.
- So if `wait` is not called, which process finishes first?
- Either one could finish first.

Zombies

- If the parent finishes first, the child becomes an orphan and is *adopted* by a system process called init whose pid is 1.
- If the child finishes first, it becomes a *zombie*.
- The child is mostly dead, but the parent might call waitpid. So its termination information must be retained until the parent either terminates or calls waitpid.

Summary

■ Basic functions

- fork spawns new processes
- exit terminates own process
- wait and waitpid wait for and reap terminated children
- execve runs new program in existing process

Lab 1 (22/1) preparation

■ Pre-lab task

- before you come to the lab session
- Setup Linux on your laptop
 - install linux in a VM
 - install Windows Subsystem for Linux (WSL)
- You will lose points if you come to the lab without finishing this task

■ In-lab task

- follow instructions